

HPE CTY WEEKLY PROGRESS DOCUMENTATION

CAPACITY ADVISOR FOR MONITORING RESOURCES

Week 1: Setting Up Apache Kafka and Prometheus on AWS EC2 (Linux)

- Installed and configured **Apache Kafka**, a distributed event streaming platform designed for high-throughput, fault-tolerant real-time data processing.
- Installed and configured **Prometheus**, an open-source monitoring and alerting toolkit optimized for time-series data and system performance metrics.

Steps followed:

APACHE KAFKA

Pre-requisites

- Launch an EC2 instance and allow port 22 for SSH.
- Launch Putty and connect to the AWS EC2 instance:
 - Install PuTTY (if not installed)
 - Create the EC2 instance with a key pair (.pem file).
 - Convert .pem to .ppk (Using PuTTYgen)
 - Connect to AWS EC2 using PuTTY:
 - Open PuTTY.
 - In Host Name, enter – ec2-user@your-aws-public-ip
 - In Connection > SSH > Auth, load the .ppk file under "Private key file for authentication".
 - Click Open.
 - Click "Yes" if you get a security alert.
- Update Your System:
 - `sudo yum update -y`
- Install Java (Required for Kafka):
 - `sudo amazon-linux-extras enable corretto8`
 - `sudo yum install java-11-amazon-corretto -y`

Downloading and Installing Kafka

- Navigate to the /opt directory:
 - `cd /opt`
- Download Kafka:
 - `wget https://downloads.apache.org/kafka/3.6.1/kafka_2.13-3.6.1.tgz`
- Extract Kafka:
 - `sudo tar -xvzf kafka_2.13-3.6.1.tgz`
 - `sudo mv kafka_2.13-3.6.1 kafka`
- Change ownership and permissions:
 - `sudo chown -R ec2-user:ec2-user /opt/kafka`

Start Zookeeper (Required for Kafka)

- Start Zookeeper:
 - `cd /opt/kafka`
 - `bin/zookeeper-server-start.sh config/zookeeper.properties`

- Run Zookeeper in the background:
 - `nohup bin/zookeeper-server-start.sh config/zookeeper.properties > zookeeper.log`

Start Kafka Broker

- Start Kafka:
 - `bin/kafka-server-start.sh config/server.properties`
- Run Kafka in the background:
 - `nohup bin/kafka-server-start.sh config/server.properties > kafka.log 2>&1 &`

Testing Kafka

- Create a Topic:
 - `bin/kafka-topics.sh --create --topic test-topic --bootstrap-server localhost:9092`
 - `--partitions 1 --replication-factor 1`
- Start a Kafka Producer to send messages
 - `bin/kafka-console-producer.sh --topic test-topic --bootstrap-server localhost:9092`
- Start a Kafka Consumer to read messages
 - `bin/kafka-console-consumer.sh --topic test-topic --from-beginning --bootstrap-server localhost:9092`

```

ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-topics.sh --create --topic test-topic --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
Created topic test-topic.
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-topics.sh --list --bootstrap-server localhost:9092
test-topic
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-console-producer.sh --topic test-topic --bootstrap-server localhost:9092
Hello
How are you?
Hi
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-topics.sh --list --bootstrap-server localhost:9092
test-topic
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-console-producer.sh --broker-list localhost:9092 --topic my-topic
How are you ?
Hello
My name
ubuntu@ip-172-31-8-99:/opt/kafka$
  
```


```

ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic my-topic --from-beginning
[2025-03-12 23:16:05,062] WARN [Consumer clientId=console-consumer, groupId=console-consumer-58065] The metadata response from the cluster reported a recoverable issue with correlation id 2 : (my-topic=LEADER_NOT_AVAILABLE) (org.apache.kafka.clients.NetworkClient)
bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic my-topic --from-beginning
^CProcessed a total of 0 messages
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic my-topic --from-beginning
^CProcessed a total of 0 messages
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic my-topic --from-beginning
^CProcessed a total of 0 messages
ubuntu@ip-172-31-8-99:/opt/kafka$ bin/kafka-console-producer.sh --broker-list localhost:9092 --topic my-topic
How are you ?
Hello
My name
  
```

PROMETHEUS

- Launch an EC2 Instance and allow port 22 for SSH, port 9090 for Prometheus, and port 9100 for Node Exporter.
- Update system packages:
 - `sudo apt update && sudo apt upgrade -y`
- Download Prometheus: `wget https://github.com/prometheus/prometheus/releases/download/v3.2.1/prometheus-3.2.1.linux-amd64.tar.gz`
- Extract and move files:
 - `tar -xvzf prometheus-3.2.1.linux-amd64.tar.gz`
 - `sudo mv prometheus-3.2.1.linux-amd64 /etc/prometheus`
- Create Prometheus user and set permissions:
 - `sudo useradd --no-create-home --shell /bin/false prometheus`

- sudo chown -R prometheus:prometheus /etc/prometheus
- Configure prometheus.yml:
 - sudo nano /etc/prometheus/prometheus.yml
- Create a systemd service for Prometheus:
 - sudo nano /etc/systemd/system/prometheus.service
- Reload systemd and start Prometheus:
 - sudo systemctl daemon-reload
 - sudo systemctl start prometheus
 - sudo systemctl enable prometheus
- Check if Prometheus is running:
 - sudo systemctl status prometheus



Problems Faced(Research) :

1. Problem in creating kafka topic

```

ec2-user@ip-172-31-11-10:~/kafka-2.13-3.9.0
├── bin
├── config
├── libs
├── logs
├── tmp
└── v
  └── A newer version of Amazon Linux is available!
      Amazon Linux 2023, GA and supported until 2028-03-15.
      https://aws.amazon.com/linux/amazon-linux-2023/

[ec2-user@ip-172-31-11-10 ~]$ cd kafka-2.13-3.9.0
[ec2-user@ip-172-31-11-10 kafka-2.13-3.9.0]$ bin/kafka-topics.sh --create --topic demo_testing2 --bootstrap-server 13.201.166.67:9092 --replication-factor 1
--partitions 1
OpenJDK 64-Bit Server VM Warning: If the number of processors is expected to increase from one, then you should configure the number of parallel GC threads appropriately using -XX:ParallelGCThreads=N
Error while executing topic command : Cannot parse argument '1--partitions' of option replication-factor
[2025-03-17 17:40:30.284] ERROR joptsimple.internal.OptionArgumentConversionException: Cannot parse argument '1--partitions' of option replication-factor
    at joptsimple.AbstractOptionsSpec.convertWith(AbstractOptionsSpec.java:92)
    at joptsimple.ArgumentAcceptingOptionSpec.convert(ArgumentAcceptingOptionSpec.java:277)
    at joptsimple.OptionSet.valueOf(OptionSet.java:223)
    at joptsimple.OptionSet.valueOf(OptionSet.java:172)
    at org.apache.kafka.tools.TopicCommand$TopicCommandOptions.valueAsOption(TopicCommand.java:852)
    at org.apache.kafka.tools.TopicCommand$TopicCommandOptions.valueAsOption(TopicCommand.java:843)
    at org.apache.kafka.tools.TopicCommand$TopicCommandOptions.replicationFactor(TopicCommand.java:907)
    at org.apache.kafka.tools.TopicCommand$CommandTopicPartition.<init>(TopicCommand.java:265)
    at org.apache.kafka.tools.TopicCommand$TopicService.createTopic(TopicCommand.java:460)
    at org.apache.kafka.tools.TopicCommand.execute(TopicCommand.java:105)
    at org.apache.kafka.tools.TopicCommand.mainNoExit(TopicCommand.java:90)
    at org.apache.kafka.tools.TopicCommand.main(TopicCommand.java:85)
Caused by: java.lang.NumberFormatException: For input string: "1--partitions"
    at joptsimple.internal.Reflection.invoke(Reflection.java:136)
    at joptsimple.internal.Reflection.invoke(Reflection.java:118)
    at joptsimple.internal.MethodInvokingValueConverter.convert(MethodInvokingValueConverter.java:48)
    at joptsimple.internal.Reflection.convertWith(Reflection.java:124)
    at joptsimple.AbstractOptionsSpec.convertWith(AbstractOptionsSpec.java:90)
    ... 11 more
Caused by: java.lang.NumberFormatException: For input string: "1--partitions"
    at java.lang.NumberFormatException.forInputString(NumberFormatException.java:65)
    at java.lang.Integer.parseInt(Integer.java:590)
    at java.lang.Integer.valueOf(Integer.java:766)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at joptsimple.internal.Reflection.invoke(Reflection.java:116)
    ... 14 more
(org.apache.kafka.tools.TopicCommand)
[ec2-user@ip-172-31-11-10 kafka-2.13-3.9.0]$

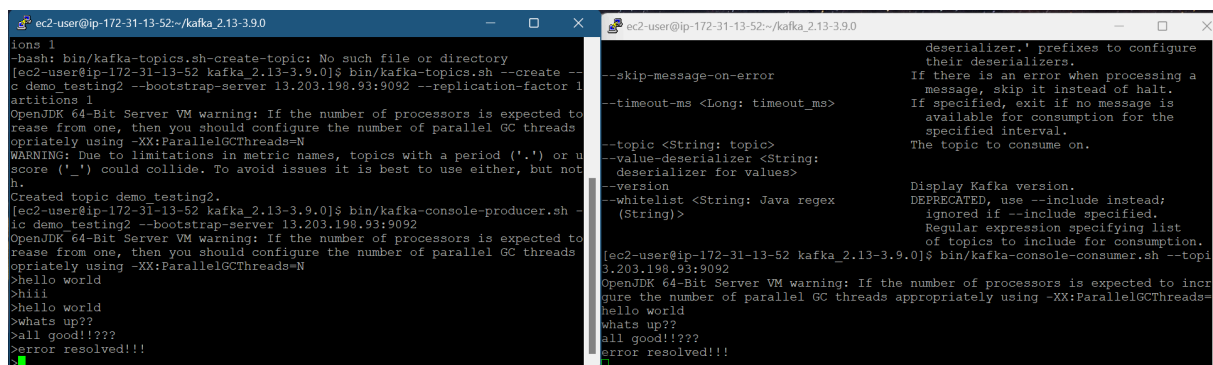
```

Troubleshoot:

Kafka topic creation issue was encountered because of insufficient memory. Therefore, we fixed it by adding swap space. Since Kafka needs more memory, we created a 2GB swap file:

- Create a 2GB swap file:
 - `sudo fallocate -l 2G /swapfile`
- Set the correct permissions:
 - `sudo chmod 600 /swapfile`
- Set up the swap space:
 - `sudo mkswap /swapfile`
- Enable the swap:
 - `sudo swapon /swapfile`

Output achieved after troubleshoot:



```
ions 1
-bash: bin/kafka-topics.sh-create-topic: No such file or directory
[ec2-user@ip-172-31-13-52 kafka_2.13-3.9.0]$ bin/kafka-topics.sh --create --
c demo_testing2 --bootstrap-server 13.203.198.93:9092 --replication-factor 1
artitions 1
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
rease from one, then you should configure the number of parallel GC threads
ppriately using -XX:ParallelGCThreads=N
WARNING: Due to limitations in metric names, topics with a period ('.') or u
score ('_') could collide. To avoid issues it is best to use either, but not
h.
Created topic demo_testing2.
[ec2-user@ip-172-31-13-52 kafka_2.13-3.9.0]$ bin/kafka-console-producer.sh --
ic demo_testing2 --bootstrap-server 13.203.198.93:9092
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
rease from one, then you should configure the number of parallel GC threads
ppriately using -XX:ParallelGCThreads=N
>hello world
>hihi
>hello world
>whats up??
>all good!!!!??
>error resolved!!!
<

deserializer.' prefixes to configure
their deserializers.
If there is an error when processing a
message, skip it instead of halt.
If specified, exit if no message is
available for consumption for the
specified interval.
The topic to consume on.

--skip-message-on-error
--timeout-ms <Long: timeout_ms>

--topic <String: topic>
--value-deserializer <String:
deserializer for values>
--version
--whitelist <String: Java regex
(String)>

Display Kafka version.
DEPRECATED, use --include instead;
ignored if --include specified.
Regular expression specifying list
of topics to include for consumption.

[ec2-user@ip-172-31-13-52 kafka_2.13-3.9.0]$ bin/kafka-console-consumer.sh --topi
3.203.198.93:9092
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to incr
gure the number of parallel GC threads appropriately using -XX:ParallelGCThreads=
hello world
whats up??
all good!!!!??
error resolved!!!
<
```