

WEEK 3 : Configuring to see the no of messages received within a time period

STEPS:

Kafka Setup:

1. Configure Zookeeper server in your VM.
2. Configure Kafka broker server in your VM.
3. Create a topic is not created before using the following command.

```
bin/kafka-topics.sh --create --topic demo_testing2 --bootstrap-server 51.20.105.68:9092 --replication-factor 1 --partitions 1
```

Node Exporter Setup:

1. Download Node exporter in your local machine or in a EC2 Instance VM.

```
Wget https://github.com/prometheus/node\_exporter/releases/download/v1.9.0/node\_exporter-1.9.0.linux-amd64.tar.gz
```

2. Unzip it using TAR.

```
tar xvfz node_exporter-1.9.0.linux-amd64.tar.gz
```

3. Run Node exporter

```
cd node_exporter-1.9.0.linux-amd64  
./node_exporter
```

NOTE:

It requires LINUX as a subsystem if you want to expose your local machine metrics.

By default, Node Exporter listens on port 9100.

1. **Verify Metrics Exposure:**
 - Open a browser and navigate to <http://localhost:9100/metrics> to see the metrics being exposed.
2. If you are running it in your local machine, then you would need to expose the port 9100 using NGROK:

If you are running node exporter in a VM , remember the endpoint where the metrics is getting shipped:

Format : <YOUR_EC2_Instance_PUBLIC_IP>/metrics

Steps:

1. Sudo apt install NGROK.
2. Configure NGROK key after registering from their website.
3. Execute Command :

```
ngrok http 9100
```

NGROK will give a endpoint for the port.

```
ngrok
  Protect endpoints w/ IP Intelligence: https://ngrok.com/r/ipintel
Session Status      online
Account             Abhilash Sarangi (Plan: Free)
Version             3.22.0
Region              India (in)
Web Interface        http://127.0.0.1:4040
Forwarding           https://3e9f-128-185-112-58.ngrok-free.app -> http://localhost:9100

Connections      ttl    opn    rt1    rt5    p50    p90
                  0      0      0.00   0.00   0.00   0.00
```

4. Note down the endpoint as we will use it in our kafka consumer script to fetch the metrics from.

Kafka Producer Script:

1. Open a google collab sheet.
2. Download kafka-python

```
!pip install kafka-python
```

```
1 !pip install kafka-python
2
Collecting kafka-python
  Downloading kafka_python-2.1.4-py2.py3-none-any.whl.metadata (9.1 kB)
  Downloading kafka_python-2.1.4-py2.py3-none-any.whl (276 kB)
    276.6/276.6 kB 4.1 MB/s eta 0:00:00
Installing collected packages: kafka-python
Successfully installed kafka-python-2.1.4
```

```

from kafka import KafkaConsumer
import json

#Kafka consumer
consumer = KafkaConsumer(
    'demo_testing2', # Topic name
    bootstrap_servers='51.20.105.68:9092', # Kafka broker (EC2 public IP)
    auto_offset_reset='earliest',
    enable_auto_commit=True,
    value_deserializer=lambda v: v.decode('utf-8') # Decode as plain text
)

print("Listening for messages on topic 'demo_testing2'...")

for message in consumer:
    try:
        json_message = json.loads(message.value)
        print(f"Received JSON message: {json_message}")
    except json.JSONDecodeError:
        print(f"Received non-JSON message: {message.value}")

```

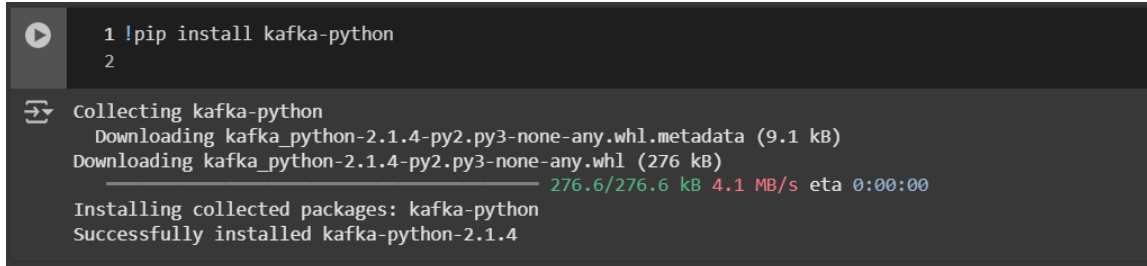
KafkaConsumer setup:

- Listens to topic demo_testing2.
- Connects to a Kafka broker running on the specified IP (51.20.105.68) and port 9092.
- auto_offset_reset='earliest': Starts reading from the beginning of the topic if no offset is stored.
- enable_auto_commit=True: Automatically commits offsets, marking messages as read.
- value_deserializer=lambda v: v.decode('utf-8'): Converts incoming byte data into readable strings.

Kafka Producer Script:

1. Open another google collab sheet.
2. Download kafka-python

```
!pip install kafka-python
```



A terminal window showing the command `!pip install kafka-python` being executed. The output shows the collection and downloading of the `kafka-python` package, including the version `2.1.4` and the file size `276.6/276.6 kB`. The installation is successful.

```
1 !pip install kafka-python
2
Collecting kafka-python
  Downloading kafka_python-2.1.4-py2.py3-none-any.whl.metadata (9.1 kB)
  Downloading kafka_python-2.1.4-py2.py3-none-any.whl (276 kB)
    276.6/276.6 kB 4.1 MB/s eta 0:00:00
Installing collected packages: kafka-python
Successfully installed kafka-python-2.1.4
```

3. Run the following script in another code block.

```
import time
from kafka import KafkaProducer
import requests
import json

# Kafka Producer Configuration
producer = KafkaProducer(
    bootstrap_servers='51.20.105.68:9092', # Replace with your Kafka
    value_serializer=lambda v: json.dumps(v).encode('utf-8') #
    # Serialize dictionary to JSON
)

while True:
    # Fetch Node Exporter Metrics
    response = requests.get('https://12e2-128-185-112-58.ngrok-
free.app/metrics') # Replace with Node Exporter URL
    metrics = response.text

    # Send a subset of metrics to Kafka
    for line in metrics.splitlines():
        if line.startswith('node_cpu_seconds_total'):
            data = {"metric": line} # Wrap metric in a dictionary
            producer.send('node_metrics', value=data)

    print("Metrics sent to Kafka topic 'node_metrics'")
    producer.flush()

    # Add delay to control frequency of fetching metrics
    time.sleep(10) # Fetch metrics every 10 seconds
```

Explanation:

- **Kafka Broker:**
 1. Ensure the IP and port (51.20.105.68:9092) are accessible.
 2. Make sure security groups/firewall rules allow traffic.
- **Serialization:**
 1. `value_serializer` uses `json.dumps` to encode Python dicts into JSON (Kafka messages must be byte-encoded).

Metrics Fetching

- **Endpoint:**
 1. Metrics are pulled from a Prometheus Node Exporter exposed via Ngrok.
 2. URL: `https://12e2-128-185-112-58.ngrok-free.app/metrics` (replace with your actual endpoint in production).
- **Parsing:**
 1. Only metrics starting with `node_cpu_seconds_total` are processed and sent.
 2. Good for reducing noise from other metrics.

Kafka Publishing

- **Topic:** Messages are sent to Kafka topic `'node_metrics'`.
- **Structure:** Each message is a JSON object like:
`{"metric": "node_cpu_seconds_total{...} 12345.0"}`
- **Flush:** Ensures messages are not stuck in buffer before the next loop.

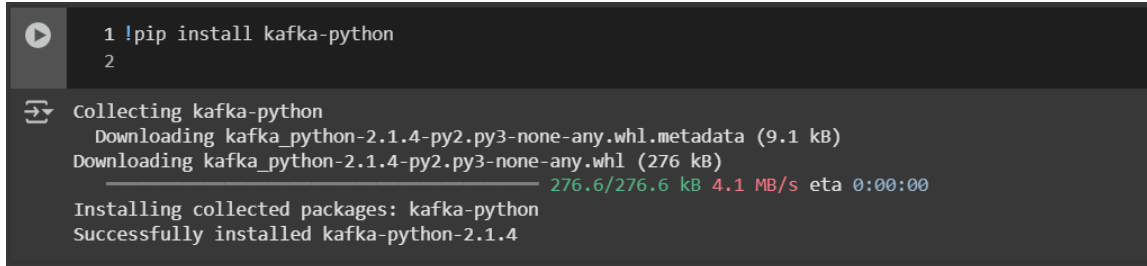
Timing

- **Loop Interval:**
 1. Metrics are fetched and sent every 10 seconds using `time.sleep(10)`.
 2. You can tune this depending on your monitoring frequency needs.

Confluent Kafka Script to fetch the no of messages sent within a time period

1. Open another google collab sheet.
2. Download kafka-python

```
!pip install kafka-python
```



A terminal window showing the command `!pip install kafka-python` and its output. The output indicates that `kafka-python` is being collected, followed by downloading the `kafka_python-2.1.4-py2.py3-none-any.whl` file (276 kB) at a speed of 4.1 MB/s. It then shows the installation of the collected packages and a successful installation of `kafka-python-2.1.4`.

```
1 !pip install kafka-python
2
Collecting kafka-python
  Downloading kafka_python-2.1.4-py2.py3-none-any.whl (276 kB)
    276.6/276.6 kB 4.1 MB/s eta 0:00:00
Installing collected packages: kafka-python
Successfully installed kafka-python-2.1.4
```

3. Run the following script

```
from confluent_kafka import Consumer, TopicPartition
from datetime import datetime, timedelta

# Kafka Configuration
KAFKA_BROKER = '51.20.105.68:9092'
TOPIC = 'node_metrics'
GROUP_ID = 'message-counter-group'

def get_message_count_last_minute():
    # Create a Kafka consumer
    consumer = Consumer({
        'bootstrap.servers': KAFKA_BROKER,
        'group.id': GROUP_ID,
        'auto.offset.reset': 'earliest',
        'enable.auto.commit': False
    })

    # Fetch metadata to get the partitions for the topic
    metadata = consumer.list_topics(timeout=10)
    if TOPIC not in metadata.topics:
        print(f"Topic '{TOPIC}' does not exist.")
        consumer.close()
        return

    partitions = metadata.topics[TOPIC].partitions.keys()

    # Get the timestamp for 1 minute ago
    one_minute_ago = int((datetime.now() -
        timedelta(minutes=1)).timestamp() * 1000)

    # Create TopicPartition objects with the timestamp
    topic_partitions = [TopicPartition(TOPIC, p, one_minute_ago) for
        p in partitions]
```

```

# Fetch offsets for the timestamp (1 minute ago)
offsets_for_time = consumer.offsets_for_times(topic_partitions)

# Calculate the number of messages in the last minute
total_messages = 0
for tp in topic_partitions:
    # Find the offset for the timestamp
    offset_info = next((o for o in offsets_for_time if o.topic
== tp.topic and o.partition == tp.partition), None)
    if offset_info is not None and offset_info.offset != -1:
        # Get the high watermark (latest offset)
        low, high =
consumer.get_watermark_offsets(TopicPartition(TOPIC, tp.partition))
        total_messages += high - offset_info.offset

consumer.close()
return total_messages

if __name__ == '__main__':
    count = get_message_count_last_minute()
    if count is not None:
        print(f"Messages received in the last minute: {count}")

```

Explanation:

This script connects to a Kafka broker and counts how many messages were produced to the `node_metrics` topic in the **last minute**.

How it works:

1. **Connects to Kafka** and checks the topic exists.
2. Calculates the **timestamp for 1 minute ago**.
3. For each partition:
 - Gets the **offset** closest to that timestamp.
 - Gets the **latest offset** (high watermark).
 - Subtracts the two to find how many messages were added since.
4. **Sums up** messages from all partitions.

Finally, it prints the total message count in the last minute.